

Università di Verona

A.Y. 2021-22

# Machine Learning & Artificial Intelligence

## Hidden Markov Models

Vittorio Murino

# Summary


1. Markov processes and models;
2. Hidden Markov Model (HMM);
3. Research and applications on HMM.

# Time series analysis

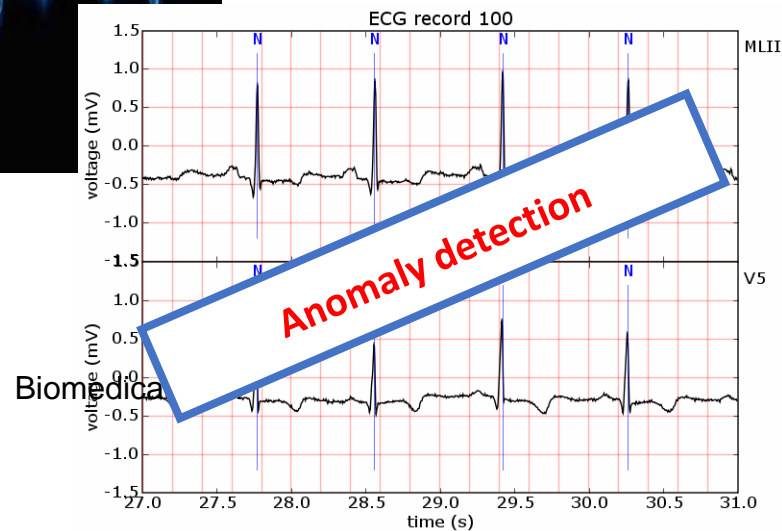
Audio signals




Financial series



Gestures



# Markov process (order 1)




$N=3$

$t=1$

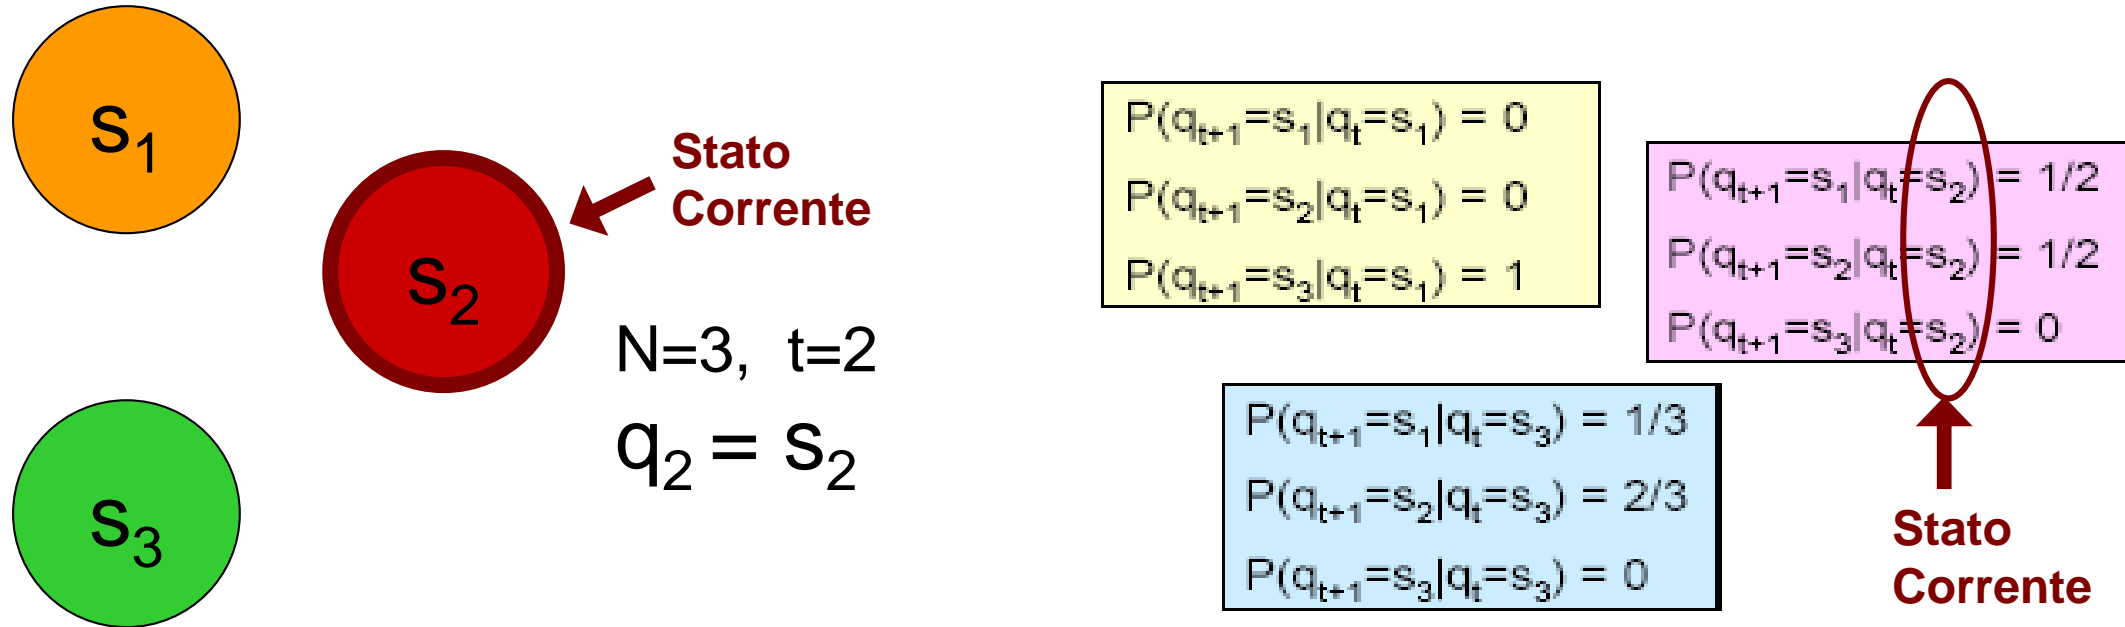
- Has  $N$  states ,  $s_1, s_2, \dots, s_N$
- It is characterized by discrete steps,  $t=1, t=2, \dots$
- The probability of starting from a certain state is dictated by the distribution:
- $=\{\pi_i\}$  :  $\pi_i = P(q_1 = s_i)$  with
$$1 \leq i \leq N, \quad \pi_i \geq 0 \quad \text{and} \quad \sum_{i=1}^N \pi_i = 1$$

# Markov process



- At the  $t$ -th instant the process is exactly in one of the available states, indicated by the variable  $q_t$
- Note:  $q_t \in \{s_1, s_2, \dots, s_N\}$
- At each iteration, the next state is chosen with a certain probability
-

# Markov process




- This probability is only determined by the previous state (*first-order markovianity*):
- $P(q_{t+1}=s_j|q_t=s_i, q_{t-1}=s_k, \dots, q_1=s_l) = P(q_{t+1}=s_j|q_t=s_i)$

# Markov hypothesis

The probability of moving to a given state depends only on the current state.

$$P(q_t = S^* | q_{t-1}, \dots, q_1) = P(q_t = S^* | q_{t-1})$$

# Markov process



$N=3, t=2$

$q_2 = s_2$

- Defining:

$$a_{i,j} = P(q_{t+1} = s_j \mid q_t = s_i)$$

I get the matrix  $N \times N$

*A transition between states,  
invariant over time:*

$A =$

$a_{1,1}$	$a_{1,2}$	$a_{1,3}$
$a_{2,1}$	$a_{2,2}$	$a_{2,3}$
$a_{3,1}$	$a_{3,1}$	$a_{3,3}$



# Markov Models


- A set of  $N$  **states**  $\mathcal{S} = \{S_1, S_2, \dots, S_N\}$
- A sequence of states  $\mathcal{Q} = \{q_1, q_2, \dots, q_T\}$

- A **transition probability matrix**

$$A = \{a_{ij} = P(q_t = S_j | q_{t-1} = S_i)\}$$

- An **initial probability distribution** over states

$$\Pi = \{\pi_i = P(q_1 = S_i)\}$$



# Markov Models

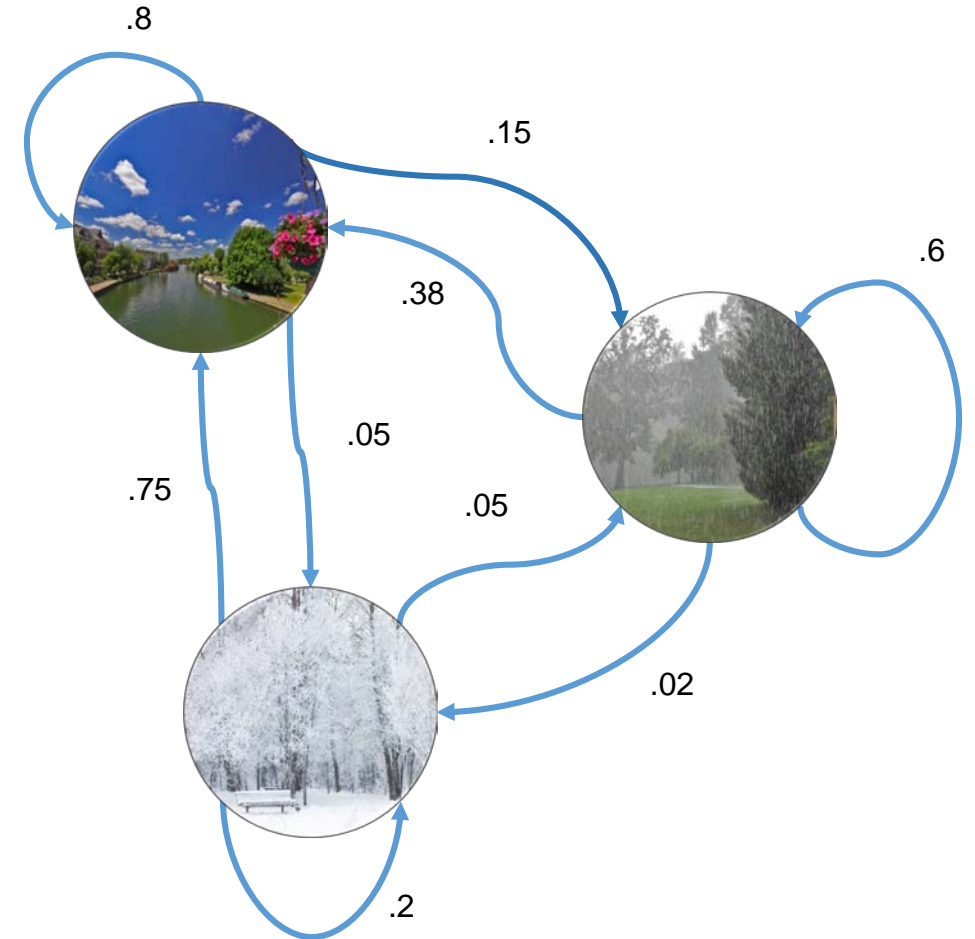
- **States:**  $\{sunny, rainy, snowy\}$

- **Transition probability matrix:**

$$A = \begin{bmatrix} .8 & .15 & .05 \\ .38 & .6 & .02 \\ .75 & .05 & .2 \end{bmatrix}$$

- **Initial probability distribution:**

$$\Pi = [.7 \quad .25 \quad .05]$$



# Exercise

- Compute the probability for the sequence:



$$P = P(Su) * P(R|Su) * P(R|R) * P(R|R) * P(Sn|R) * P(Sn|Sn)$$

$$P = 0.0001512$$

# Features of Markovian processes

- They are (discrete) processes characterized by:
  - Markovianity of the first order
  - stationary
  - have an initial distribution
- Knowing the above features, one can exhibit a **(probabilistic) model of Markov (MM)** as

$$\lambda = (A, \pi)$$

# What is a stochastic or probabilistic model for?

- Models and reproduces **stochastic processes**
- Describes by probability the **causes that lead from one state of the system to another**
- In other words, the more likely you are to move from state A to state B, the more likely it is that **A will cause B**

# What can be done on a probabilistic model?


- **Training**

- The constituent elements of the model are estimated

- **Inferences of various kinds (I question the model):**

- Probability of a sequence of states, given the model
  - Invariant properties etc.

# What is a Markov model for?

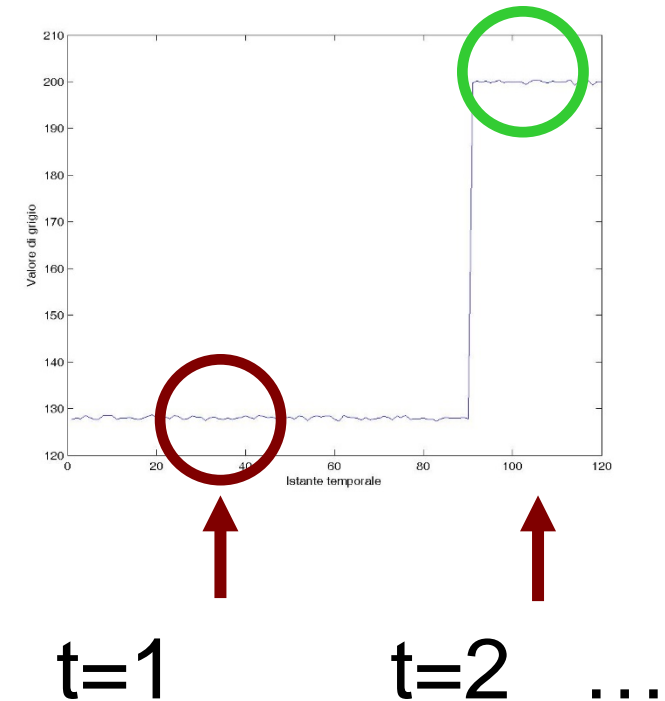


- Model Markovian *stochastic behaviors* (of order  $N$ ) of a system in which the states are:
  - **Explicit** (I can give them a name)
  - **Observable** (I have observations that uniquely identify the state)

# Example: Traffic light





- It is a system whose states are:
  - **Explicit** (the different lamps lit)
  - **Observable** (the colors of the lamps I observe)







# Traffic light – trained model



$$\pi = \begin{array}{|c|c|c|} \hline \pi_1 = 0.33 & \pi_2 = 0.33 & \pi_3 = 0.33 \\ \hline \end{array}$$

$$A = \begin{array}{|c|c|c|} \hline a_{11} = 0.1 & a_{12} = 0.9 & a_{13} = 0 \\ \hline a_{21} = 0.01 & a_{22} = 0 & a_{23} = 0.99 \\ \hline a_{31} = 1 & a_{32} = 0 & a_{33} = 0 \\ \hline \end{array}$$

# Traffic light – inference



$$O_2 = \langle q_2 = s_3, q_1 = s_2 \rangle$$

Inference:


$$P(O | \lambda) = P(O) =$$

$$= P(q_2 = s_3, q_1 = s_2) = P(q_2, q_1)$$


# Important inference

$$\begin{aligned} P(q_t, q_{t-1}, \dots, q_1) &= P(q_t | q_{t-1}, \dots, q_1) P(q_{t-1}, \dots, q_1) \\ &= P(q_t | q_{t-1}) P(q_{t-1}, q_{t-2}, \dots, q_1) \\ &= P(q_t | q_{t-1}) P(q_{t-1} | q_{t-2}) P(q_{t-2}, \dots, q_1) \\ &\dots \\ &= P(q_t | q_{t-1}) P(q_{t-1} | q_{t-2}) \dots P(q_2 | q_1) P(q_1) \end{aligned}$$



# Graphic representation



The graphic structure of such a joint probability is written in this form, where



# Traffic light – inferences, response



$$P(O | \lambda) = P(O)$$

$$= P(q_2 = s_3, q_1 = s_2)$$

$$= P(q_2 = s_3 | q_1 = s_2) P(q_1 = s_2)$$

$$= 0.99 * 0.33 = 0.326$$

## Second important inference

- Probability calculation  $P(q_T = s_j)$

- STEP 1:

I evaluate how to calculate  $P(Q)$  for each path of states

$$Q = \{q_1, q_2, \dots, q_T = s_j\}, \text{ that is}$$

$$P(q_T, q_{T-1}, \dots, q_1)$$

- STEP 2:

I use this method to calculate  $P(q_T = s_j)$ , that is:

$$\circ P(q_T = s_j) = \sum P(Q)$$

$Q \in \text{paths of length } T \text{ ending in } s_j$

Onerous calculation: EXPONENTIAL in  $T$  ( $O(N^T)$ )!

## Second important inference (2)

- **Idea:** for each state  $s_j$  we call  $p_T(j)$  = prob. to be in the state  $s_j$  at the time  $T \rightarrow P(q_T = s_j)$ ;
  - It can be defined by induction:

$$\forall i \quad p_1(i) = \begin{cases} 1 & \text{if } s_i \text{ is the current state} \\ 0 & \text{otherwise} \end{cases}$$

$$\forall j \quad p_{t+1}(j) = P(q_{t+1} = s_j) = \sum_{i=1}^N P(q_{t+1} = s_j, q_t = s_i)$$

## Second important inference (3)



$$\begin{aligned} p_{t+1}(j) &= \sum_{i=1}^N P(q_{t+1} = s_j, q_t = s_i) = \\ &= \sum_{i=1}^N P(q_{t+1} = s_j \mid q_t = s_i) P(q_t = s_i) = \sum_{i=1}^N a_{ij} p_t(i) \end{aligned}$$

- I use this method starting from  $P(q_T = s_j)$  and proceeding backwards
- The cost of computation in this case is  $O(TN^2)$




# Limits of Markovian models


1. The state should always be ***observable deterministically***, observations have no noise.




# Limits of Markovian models



OK



NO




# Limits of Markovian models



2. **(and more important!)** In the case of the traffic light the state is **explicit**, (a particular traffic light configuration), and **can be assessed directly through observation**

(the status corresponds to the color of the traffic light)

- This is not always the case!




# Limits of Markovian models




# Limits of Markovian models



- *I watch* the video sequence:  
*I observe* that there is a **system that evolves**, but I cannot understand which the regulatory states are.




# Limits of Markovian models




- *I watch* the video sequence: *I observe* that there is a **system that evolves**, but I cannot understand which the regulatory states are.
- The system, however, evolves in **states**, which I understand by *observing* the phenomenon

# Limits of Markovian models



- Better: the system evolves thanks to **"hidden" states**, the states of the traffic light, which I do not see and I do not know even the existence.
- I don't observe the states, but I can only **observe** the probable "consequences" of such states, i.e. the flows of cars


# Limits of Markovian models



- I don't name the states, I just consider as hidden entities and *identifiable only through observations* (cars' flows)
- I can establish a **relationship between *observation* and *hidden state***



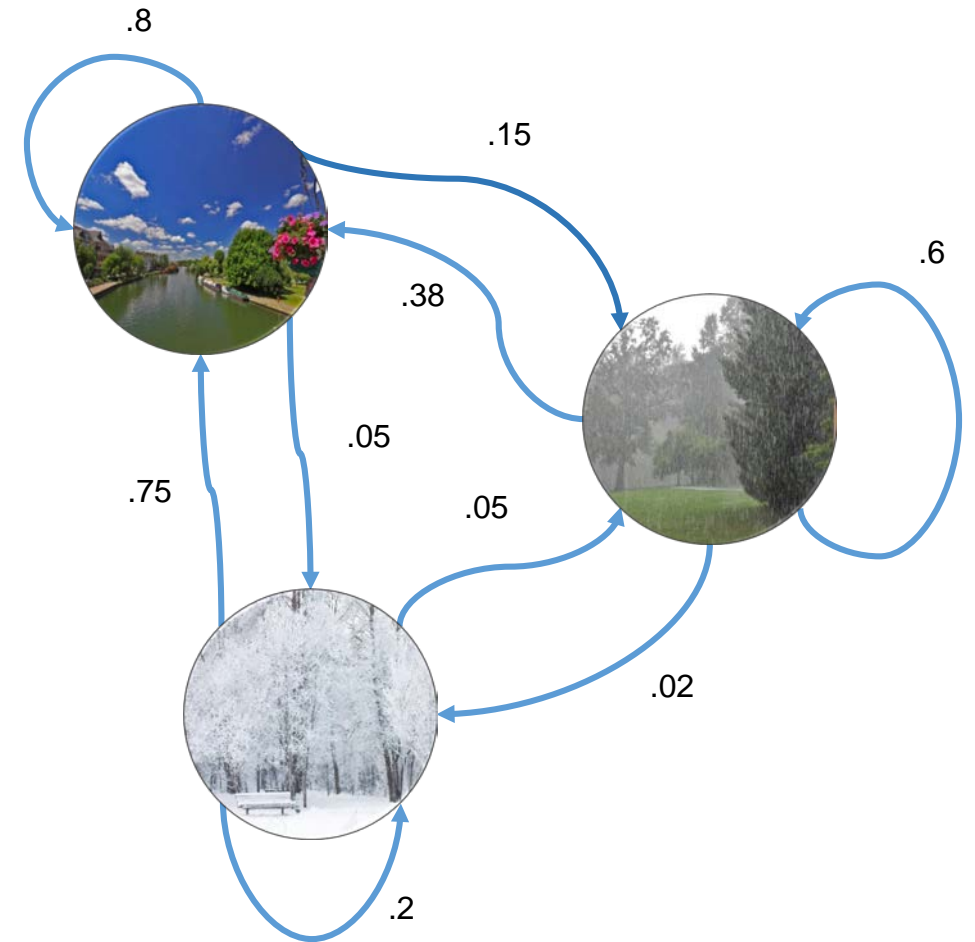
# Markov Models with Hidden states or Hidden Markov Models (HMMs)



- Hidden Markov Model fits into this context
- They **probabilistically** describe the *system dynamics* avoiding to directly identify its causes, or rather, seeking to estimate them
- States are identifiable only by *observations*, in a **probabilistic manner**.

# Hidden Markov Models

**States are not observable!**



# Hidden Markov Model (HMM)

- Statistical sequence classifier, widely used in different contexts.
- Such a model can be understood as an extension of the Markov model from which it differs for the **unobservability of its states**.
- Each state has associated a probability function that describes the probability that a certain symbol (output) is emitted by that state.

# HMM: a formal definition

- From [Rabiner 89]:

“The Hidden Markov Model is a doubly embedded stochastic process with an underlying stochastic process that is *not* observable (it is hidden), but can only be observed through another set of stochastic processes that produce the sequence of observations”


L.R. Rabiner. **A tutorial on hidden Markov models and selected applications in speech recognition.**  
*Proceedings of the IEEE*, Vol. 77, Issue 2, Feb. 1989.

# HMM: formal definition

- An HMM (discrete) consists of:
  - A set  $S = \{s_1, s_2, \dots, s_N\}$  of hidden states
  - A transition matrix  $A = \{a_{ij}\}$   
between hidden states  $1 \leq i, j \leq N$
  - An initial distribution over hidden states  $\pi = \{\pi_i\}$


$$\pi = \begin{array}{|c|c|c|} \hline \pi_1 = 0.33 & \pi_2 = 0.33 & \pi_3 = 0.33 \\ \hline \end{array}$$

$$A = \begin{array}{|c|c|c|} \hline a_{11} = 0.1 & a_{12} = 0.9 & a_{13} = 0 \\ \hline a_{21} = 0.01 & a_{22} = 0.2 & a_{23} = 0.79 \\ \hline a_{31} = 1 & a_{32} = 0 & a_{33} = 0 \\ \hline \end{array}$$



# HMM: formal definition

- A set  $V = \{v_1, v_2, \dots, v_M\}$  of observation symbols
- A probability distribution on observation symbols  $B = \{b_{jk}\}$ , which indicates the probability of emission of the symbol  $v_k$  when the system state is  $s_j$ .



$B =$

$b_{11}=0.8$	$b_{21}=0.1$	$b_{31}=0.1$
$b_{12}= 0.1$	$b_{22}= 0.8$	$b_{32}= 0.1$
$\vdots$	$\vdots$	$\vdots$
$b_{1M}= 0.1$	$b_{2M}= 0.1$	$b_{3M}=0.8$


# HMM: formal definition

- We denote an HMM with a triple  $\lambda = (A, B, \pi)$

$$\pi = \begin{array}{|c|c|c|} \hline \pi_1 = 0.33 & \pi_2 = 0.33 & \pi_3 = 0.33 \\ \hline \end{array}$$

$$A = \begin{array}{|c|c|c|} \hline a_{11} = 0.1 & a_{12} = 0.9 & a_{13} = 0 \\ \hline a_{21} = 0.01 & a_{22} = 0.2 & a_{23} = 0.79 \\ \hline a_{31} = 1 & a_{32} = 0 & a_{33} = 0 \\ \hline \end{array}$$


$$B = \begin{array}{|c|c|c|} \hline b_{11} = 0.8 & b_{21} = 0.1 & b_{31} = 0.1 \\ \hline b_{12} = 0.1 & b_{22} = 0.8 & b_{32} = 0.1 \\ \hline \vdots & \vdots & \vdots \\ \hline b_{1M} = 0.1 & b_{2M} = 0.1 & b_{3M} = 0.8 \\ \hline \end{array}$$



# Assumptions on observations

- Conditional independence


$$P(o_t = X \mid q_t = s_j, q_{t-1}, q_{t-2}, \dots, q_2, q_1, o_{t-1}, o_{t-2}, \dots, o_2, o_1) \\ = P(o_t = X \mid q_t = s_j)$$



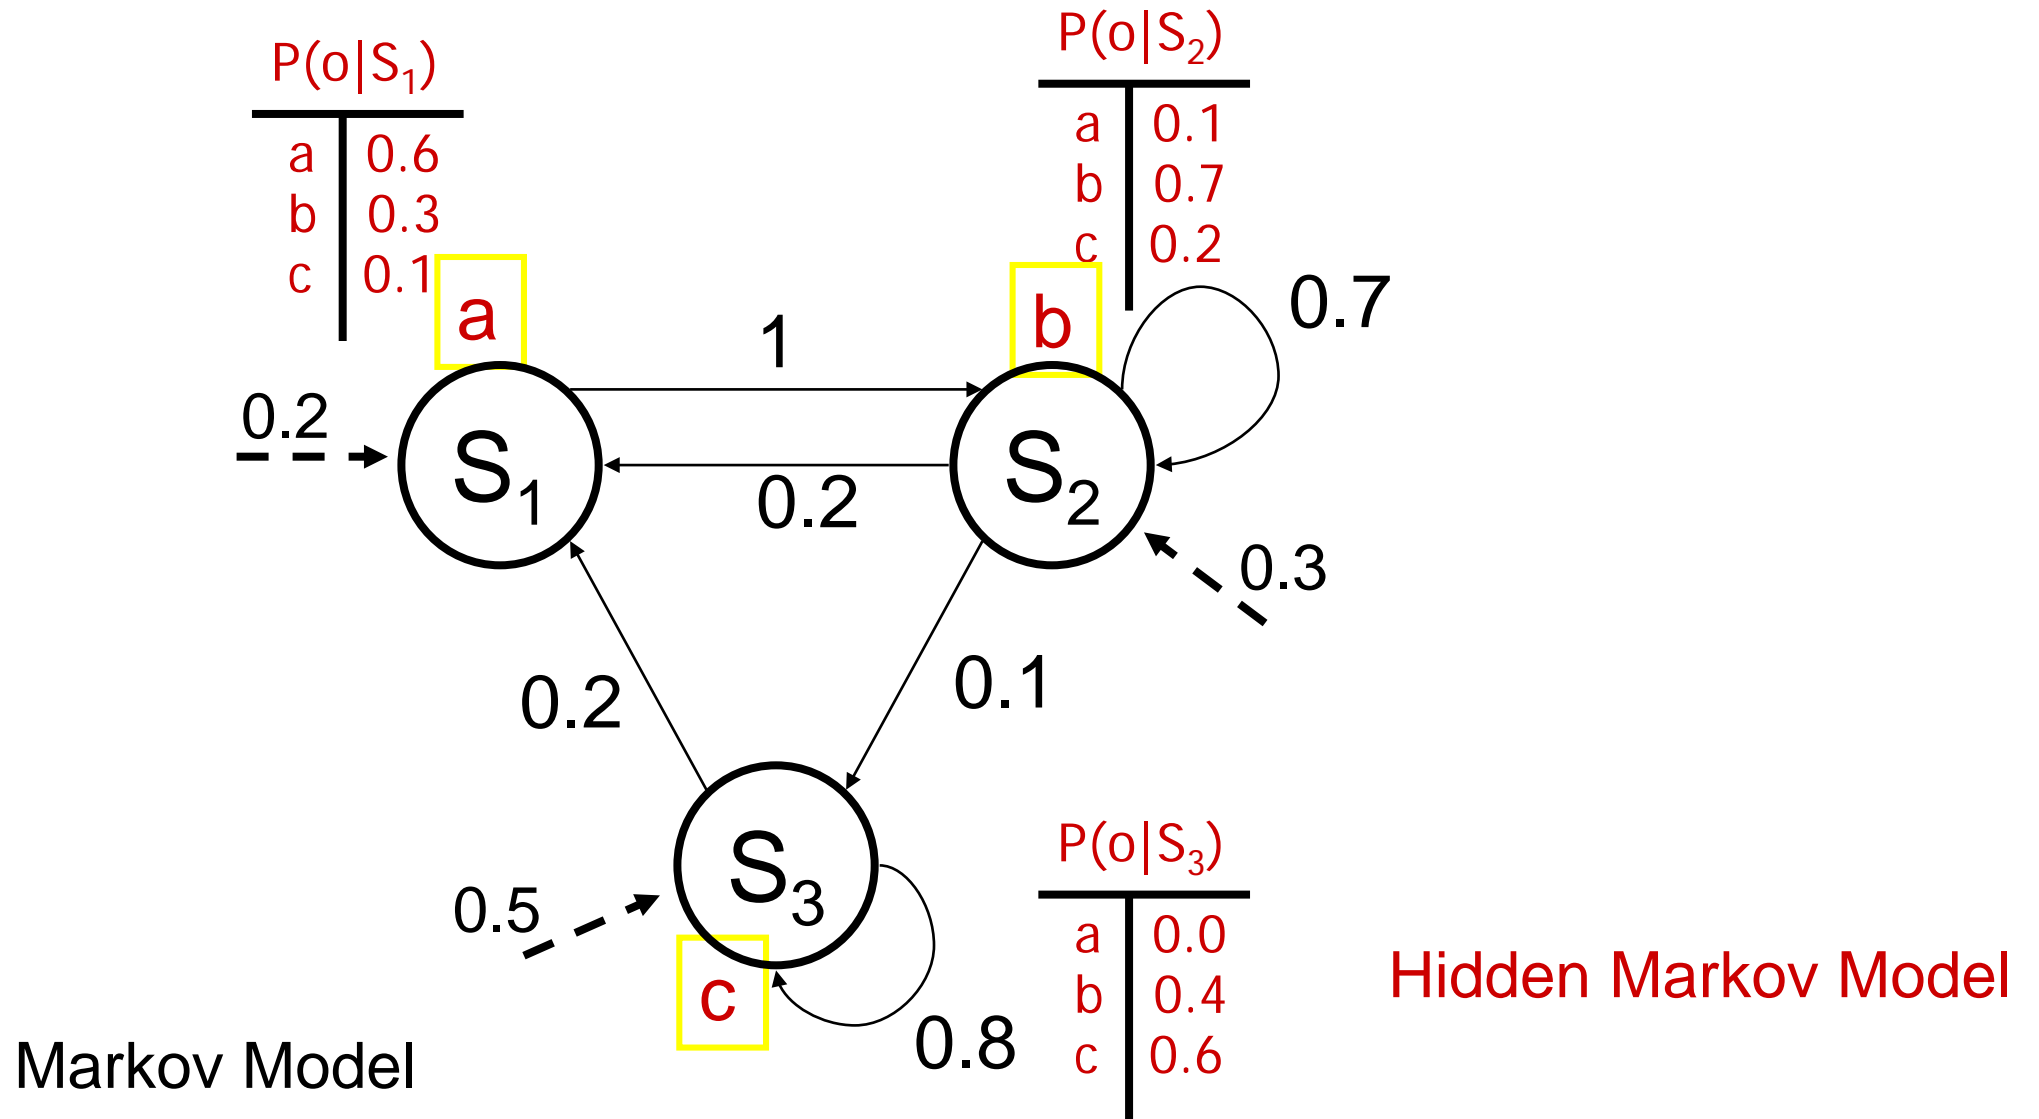


# Hidden Markov Models


- A set of  $N$  **states**  $\mathcal{S} = \{S_1, \dots, S_N\}$
- A sequence of states  $\mathcal{Q} = q_1, \dots, q_T$
- An **initial probability distribution** over states  
 $\Pi = \{\pi_i = P(q_1 = S_i)\}$
- A **transition probability matrix**  
 $A = \{a_{ij} = P(q_t = S_j | q_{t-1} = S_i)\}$
- A set of **emission probabilities**  
 $B = b_i(v_k) = P(o_t = v_k | q_t = S_i)$
- An **observation vocabulary**  $\mathcal{V} = \{v_1, \dots, v_M\}$
- A sequence of **observations**  $\mathcal{O} = o_1, \dots, o_T$



# From a Markov Model to a Hidden Markov Model



# Hidden Markov Models



# Key problems for HMMs

## Problem 1: Evaluation o Likelihood

Given an HMM  $\lambda$  model and an observation string  $\mathbf{O}=(O_1, O_2, \dots, O_t, \dots, O_T)$  calculate  $P(\mathbf{O} | \lambda)$   
→ *forward procedure*


## Problem 2: Decoding

Given an observation string  $\mathbf{O}$  and an HMM  $\lambda$  model, calculate the most likely sequence of states  $S_1 \dots S_T$  that generated  $\mathbf{O}$   
→ *Viterbi procedure*

## Problem 3: Training

Given a set of observations  $\{\mathbf{O}\}$ , determine the best HMM model  $\lambda=(\pi, A, B)$ , i.e. the model for which  $P(\mathbf{O} | \lambda)$  is maximized  
→ *Baum Welch procedure (forward-backward)*


# HMM – string generator



- 3 states:  $s_1, s_2, s_3$
- 3 symbols: O, CA, OR

$\pi_1 = 0.33$	$\pi_2 = 0.33$	$\pi_3 = 0.33$
$b_1(O) = 0.8$	$b_2(O) = 0.1$	$b_3(O) = 0.1$
$b_1(OR) = 0.1$	$b_2(OR) = 0.0$	$b_3(OR) = 0.8$
$b_1(CA) = 0.1$	$b_2(CA) = 0.9$	$b_3(CA) = 0.1$
$a_{11} = 0$	$a_{12} = 1$	$a_{13} = 0$
$a_{21} = 1/3$	$a_{22} = 2/3$	$a_{23} = 0$
$a_{31} = 1/2$	$a_{32} = 1/2$	$a_{33} = 0$

# HMM – string generator



Our problem is that the states are not directly observable!

$q_0 =$	$S_2$	$O_1 =$	CA
$q_1 =$	$S_2$	$O_2 =$	CA
$q_2 =$	$S_1$	$O_3 =$	O

$q_0 =$	?	$O_1 =$	CA
$q_1 =$	?	$O_2 =$	CA
$q_2 =$	?	$O_3 =$	O

# Problem 1: Probability of a series of observations


- $P(\mathbf{O}) = P(O_1, O_2, O_3) = P(O_1 = \text{CA}, O_2 = \text{CA}, O_3 = \text{O})?$
- Brute force strategy:

$$P(\mathbf{O}) = \sum_{\mathbf{Q} \in \text{paths of length 3}} P(\mathbf{O}, \mathbf{Q})$$

$\mathbf{Q} \in \text{paths of length 3}$

$$= \sum_{\mathbf{Q} \in \text{paths of length 3}} P(\mathbf{O} | \mathbf{Q}) P(\mathbf{Q})$$


$\mathbf{Q} \in \text{paths of length 3}$



# Problem 1: Probability of a series of observations

- $P(\mathbf{O}) = P(O_1, O_2, O_3) = P(O_1 = X, O_2 = X, O_3 = Z)$ ?
- Brute force strategy:

$$\begin{aligned} P(\mathbf{O}) &= \sum P(\mathbf{O}, \mathbf{Q}) \\ &= \sum P(\mathbf{O} | \mathbf{Q}) P(\mathbf{Q}) \end{aligned}$$



$$\begin{aligned} P(\mathbf{Q}) &= P(q_1, q_2, q_3) = \\ &= P(q_1) P(q_2, q_3 | q_1) \\ &= P(q_1) P(q_2 | q_1) P(q_3 | q_2) \end{aligned}$$

For example, in the case


$$\begin{aligned} \mathbf{Q} &= S_2 S_2 S_1 = \pi_2 a_{22} a_{21} \\ &= 1/3 * 2/3 * 1/3 = 2/27 \end{aligned}$$



# Problem 1: Probability of a series of observations

- $P(\mathbf{O}) = P(O_1, O_2, O_3) = P(O_1=X, O_2=X, O_3=Z)$ ?
- Brute force strategy:

$$\begin{aligned} \circ P(\mathbf{O}) &= \sum P(\mathbf{o}, \mathbf{Q}) \\ &= \sum P(\mathbf{O} | \mathbf{Q}) P(\mathbf{Q}) \end{aligned}$$



$$\begin{aligned} P(\mathbf{O} | \mathbf{Q}) &= \\ &= P(O_1, O_2, O_3 | q_1, q_2, q_3) \\ &= P(O_1 | q_1) P(O_2 | q_2) P(O_3 | q_3) \end{aligned}$$

For example, in the case

$$\begin{aligned} \mathbf{Q} &= S_2 S_2 S_1 = \\ &= 9/10 * 9/10 * 8/10 = 0.648 \end{aligned}$$

# Considerations

- Previous calculations solve **only one term of the summation**: for the calculation of  $P(\mathbf{O})$  are required  $27 P(Q)$  and  $27 P(\mathbf{O} | Q)$
- For a sequence of 20 observations we need  $3^{20} P(Q)$  and  $3^{20} P(\mathbf{O} | Q)$
- There is a more effective way, which is based on the definition of a particular probability
- Generally:

$$P(\mathbf{O} | \lambda) = \sum_{\text{All sequences } Q_1, \dots, Q_T} \pi_{Q_1} b_{Q_1}(O_1) a_{Q_1 Q_2} b_{Q_2}(O_2) a_{Q_2 Q_3} \dots$$

is of high complexity,  $O(N^T T)$ , where  $N$  is the number of states,  $T$  length of the sequence

# Forward Procedure

- Given the observations  $O_1, O_2, \dots, O_T$  we define

$$\alpha_t(i) = P(O_1, O_2, \dots, O_t, q_t = s_i \mid \lambda), \text{ where } 1 \leq t \leq T$$

that is:

- *we have seen the first  $t$  observations*
- *we ended up in  $s_i$  at the  $t$ -th visited state*



- This probability can be defined recursively:

$$\alpha_1(i) = P(O_1, q_1 = s_i) = P(q_1 = s_i)P(O_1 \mid q_1 = s_i) = \pi_i b_i(O_1)$$

- By inductive hypothesis  $\alpha_t(i) = P(O_1, O_2, \dots, O_t, q_t = s_i \mid \lambda)$
- I want to calculate:

$$\alpha_{t+1}(j) = P(O_1, O_2, \dots, O_t, O_{t+1}, q_{t+1} = s_j \mid \lambda)$$

$$\begin{aligned}
\alpha_{t+1}(j) &= P(O_1, O_2, \dots, O_t, O_{t+1}, q_{t+1}=s_j) \\
&= \sum_{i=1}^N P(O_1, O_2, \dots, O_t, q_t=s_i, O_{t+1}, q_{t+1}=s_j) \\
&= \sum_{i=1}^N P(O_{t+1}, q_{t+1}=s_j | O_1, O_2, \dots, O_t, q_t=s_i) P(O_1, O_2, \dots, O_t, q_t=s_i) \\
&= \sum_{i=1}^N \underbrace{P(O_{t+1}, q_{t+1}=s_j | q_t=s_i)}_{\text{p.i.i.}} \alpha_t(i) \\
&= \sum_{i=1}^N \underbrace{P(q_{t+1}=s_j | q_t=s_i)}_{a_{ij}} \underbrace{P(O_{t+1} | q_{t+1}=s_j)}_{b_j(O_{t+1})} \alpha_t(i) \\
&= \sum_{i=1}^N [a_{ij} \alpha_t(i)] b_j(O_{t+1})
\end{aligned}$$



# Response to Problem 1: Evaluation

- Given  $O_1, O_2, \dots, O_t, \dots, O_T$  and knowing  $\alpha_t(i) = P(O_1, O_2, \dots, O_t, q_t = s_i | \lambda)$ , we can calculate:

$$P(O | \lambda) = P(o_1, o_2, \dots, o_T | \lambda) = \sum_{i=1}^N \alpha_T(i)$$

having complexity  $O(N^2T)$

- But also other useful quantities, for example:

$$P(q_t = s_i | O_1, O_2, \dots, O_t) = \frac{\alpha_t(i)}{\sum_{j=1}^N \alpha_t(j)}$$

# Response to Problem 1: Evaluation

- $\alpha$  is called a *forward* variable
- Alternatively, it can be calculated recursively by introducing another variable, the so-called *backward variable*

$$\begin{aligned}\beta_t(j) &= P(O_{t+1} \dots O_T \mid q_t = s_j, \lambda) = \sum_{i=1}^N P(O_{t+1} \dots O_T, q_{t+1} = s_i \mid q_t = s_j, \lambda) \\ &= \sum_{i=1}^N \beta_{t+1}(i) a_{ji} b_i(O_{t+1})\end{aligned}$$

Hence

$$\begin{aligned}P(O \mid \lambda) &= \sum_{j=1}^N \alpha_t(j) \beta_t(j) \quad \forall t \\ &= \sum_{j=1}^N \beta_1(j) \quad \text{Please, verify!}\end{aligned}$$

## Problem 2: Decoding (more likely path)

- What is the most probable (state) path (MPP) that generated  $O_1, O_2, \dots, O_T$ ? That is, how to compute:

$$\operatorname{argmax}_{\mathbf{Q}} P(\mathbf{Q} \mid O_1 O_2 \dots O_T) ?$$

- Brute force strategy:

$$\operatorname{argmax}_{\mathbf{Q}} \frac{P(O_1 O_2 \dots O_T \mid \mathbf{Q}) P(\mathbf{Q})}{P(O_1 O_2 \dots O_T)}$$

$$\propto \operatorname{argmax}_{\mathbf{Q}} P(O_1 O_2 \dots O_T \mid \mathbf{Q}) P(\mathbf{Q})$$

# Efficient computation of the MPP

- Let's define the following probability:

$$\delta_t(i) = \max_{q_1 q_2 \dots q_{t-1}} P(q_1 q_2 \dots q_{t-1}, q_t = s_i, O_1 O_2 \dots O_t)$$

i.e., the maximum probability of paths of length  $t-1$  which:

- occur,
  - end up in the state  $s_i$  at time  $t$ ,
  - produce as output  $O_1, O_2, \dots, O_t$
- You look for the single best sequence of single states (path) maximizing  $P(\mathbf{Q}|\mathbf{O}, \lambda)$
  - The solution to this problem is a dynamic programming technique called Viterbi's algorithm.
    - We look for the most likely single state at the  $i$ -th position given the previous observations and states



# Viterbi algorithm

1) Initialization:

$$\delta_1(i) = \pi_i b_i(O_1), \quad 1 \leq i \leq N$$

$$\psi_1(i) = 0.$$

By induction we have

$$\delta_{t+1}(j) = [\max_i \delta_t(i) a_{ij}] \cdot b_j(O_{t+1}).$$

# Viterbi algorithm


2) Recursion:

$$\delta_t(j) = \max_{1 \leq i \leq N} [\delta_{t-1}(i) a_{ij}] b_j(O_t), \quad 2 \leq t \leq T$$

$$1 \leq j \leq N$$

$$\psi_t(j) = \operatorname{argmax}_{1 \leq i \leq N} [\delta_{t-1}(i) a_{ij}], \quad 2 \leq t \leq T$$

$$1 \leq j \leq N.$$



ATTENTION:  
calculated for each j !!!

# Viterbi algorithm


## 3) Termination:

$$p^* = \max_{1 \leq i \leq N} [\delta_T(i)]$$

$$q_T^* = \operatorname{argmax}_{1 \leq i \leq N} [\delta_T(i)].$$

## 4) Path (state sequence) backtracking:

$$q_t^* = \psi_{t+1}(q_{t+1}^*), \quad t = T - 1, T - 2, \dots, 1.$$



## Problem 3: HMM Training

- We talk about HMM training process or parameter estimation phase, in which the parameters of  $\lambda=(A,B, \pi)$ , are estimated from *training observations*
- Usually, the Maximum Likelihood estimate is used

$$\lambda^* = \underset{\lambda}{\operatorname{argmax}} \quad P(O_1 O_2 \dots O_T \mid \lambda)$$

- But other estimates can also be used

$$\underset{\lambda}{\operatorname{max}} \quad P(\lambda \mid O_1 O_2 \dots O_T)$$

# ML estimation of HMM:

## Baum Welch's re-estimation procedure

Let's define

- $\gamma_t(i) = P(q_t = s_i \mid O_1 O_2 \dots O_T, \lambda)$
- $\xi_t(i, j) = P(q_t = s_i, q_{t+1} = s_j \mid O_1 O_2 \dots O_T, \lambda)$

These quantities can be calculated efficiently (cf. Rabiner)

$$\sum_{j=1}^N \xi_t(i, j) = \gamma_t(i)$$

$$\sum_{t=1}^{T-1} \xi_t(i, j) = \text{expected number of transitions from state } i \text{ to state } j \text{ during the journey}$$


$$\sum_{t=1}^{T-1} \gamma_t(i) = \text{expected number of transitions passing through state } i \text{ along the way}$$

- Using forward and backward variables,  $\xi$  is also calculable as

$$\xi_t(i, j) = \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{P(O|\lambda)}$$

$$= \frac{\alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}{\sum_{i=1}^N \sum_{j=1}^N \alpha_t(i) a_{ij} b_j(O_{t+1}) \beta_{t+1}(j)}$$

(E step)



# ML estimation of HMM:

## Baum Welch's re-estimation procedure

$\bar{\pi}_i$  = expected frequency (number of times) in state  $S_i$  at time ( $t = 1$ ) =  $\gamma_1(i)$

$\bar{a}_{ij}$  =  $\frac{\text{expected number of transitions from state } S_i \text{ to state } S_j}{\text{expected number of transitions from state } S_i}$

$$= \frac{\sum_{t=1}^{T-1} \xi_t(i, j)}{\sum_{t=1}^{T-1} \gamma_t(i)}$$

$\bar{b}_j(k)$  =  $\frac{\text{expected number of times in state } j \text{ and observing symbol } v_k}{\text{expected number of times in state } j}$

$$= \frac{\sum_{t=1}^T \gamma_t(j) \text{ s.t. } O_t = v_k}{\sum_{t=1}^T \gamma_t(j)}$$

Parameter re-estimation formulas  
(M step)

# Baum-Welch algorithm

- These quantities are used in the process of estimating HMM parameters iteratively
- A variation of the Expectation-Maximization (EM) algorithm is used
  - that performs a local optimization
  - maximizing the log-likelihood of the model with respect to the data

$$\lambda_{\text{opt}} = \operatorname{argmax} \log P(\{\mathbf{O}_1\} \mid \lambda)$$



## EM - BAUM WELCH (2)

- Knowing the quantities such as:
  - expected number of transitions leaving state  $i$  along the way,
  - expected number of transitions from state  $i$  to state  $j$  along the path,
  - we could calculate the current ML estimates of  $\lambda$  ( $=\bar{\lambda}$ ), that is

$$\bar{\lambda} = (\bar{A}, \bar{B}, \bar{\pi})$$

- These considerations give rise to the Baum-Welch algorithm

- Algorithm:

- 1) I initialize the model  $\bar{\lambda} \equiv (A_0, B_0, \pi_0)$
- 2) the current model is  $\lambda = \bar{\lambda}$
- 3) I use the model  $\lambda$  to calculate the right part of the re-estimation formulas, i.e., the statistics (E step)
- 4) I use such statistics for the re-estimation of parameters obtaining the new model  $\bar{\lambda} = (\bar{A}, \bar{B}, \bar{\pi})$  (M step)
- 5) go to step 2, until termination occurs

- Baum showed that at every step:

$$P(O_1, O_2, \dots, O_T \mid \bar{\lambda}) > P(O_1, O_2, \dots, O_T \mid \lambda)$$


- Usual termination conditions:

- after a fixed number of cycles
- convergence of the likelihood value

# HMM training

Fundamental issue:

- Baum-Welch is a gradient-descent optimization technique (local optimizer)
- the log-likelihood is highly multimodal



- initialization of parameters can crucially affect the convergence of the algorithm

# Some open issues/research trends

## 1. Model selection

- how many states?
- which topology?

## 2. Extending standard models

- modifying dependencies or components

## 3. Injecting discriminative skills into HMM

# Some open issues/research trends

## 1. Model selection

- how many states?
- which topology?

## 2. Extending standard models

- modifying dependencies or components

## 3. Injecting discriminative skills into HMM

# Model selection

- The problem of determining the HMM structure:
  - not a new problem, but still a not completely solved issue
- 1. Choosing the number of states: the “standard” model selection problem
- 2. Choosing the topology: forcing the absence or the presence of connections


# Model selection problem 1: selecting the number of states

- Number of states: prevents overtraining
- The problem could be addressed using standard model selection approaches




... let's understand the concept with a toy example

# What is model selection?

Toy example: some experimental data to which we want to fit a polynomial.




The model selection question is: which order?





# What is model selection?



Model selection goal:

how to identify the underlying trend of the data, ignoring the noise?

# Model selection: solutions

- Typical solution (usable for many probabilistic models)
  - train several models with different orders  $k$
  - choose the one maximizing an “optimality” criterion


Which “optimality” criterion?

- First naive solution: maximizing likelihood of data w.r.t. model

# Maximizing Log Likelihood

- Problem: Log Likelihood is not decreasing when augmenting the order

Not applicable criterion!



## Alternative: penalized likelihood

- Idea: find a compromise between fitting accuracy and simplicity of the model
- Insert a “penalty term” which grows with the order of the model and discourages highly complex models

$$K_{\text{best}} = \operatorname{argmax}_k ( LL(y | \theta_k) - C(k) )$$

↑  
complexity penalty

Examples: BIC, MDL, MML, AIC, ...

## Alternative: penalized likelihood





- Example: Bayesian information criterion (BIC) [Schwartz, 1978]

$$k_{\text{best}} = \arg \max_k \left\{ \text{LL}(y \mid \theta_k) - \frac{k}{2} \log(n) \right\}$$

increases with  $k$


decreases with  $k$   
(penalizes larger  $k$ )

## Back to the polynomial toy example




# Model selection problem 2: selecting the best topology


- Problem: forcing the absence or the presence of connections
- Typical ad-hoc solutions
  - ergodic HMM (no constraints)
  - left to right HMM (for speech)
  - circular HMM (for shape recognition)



standard ergodic HMM



circular HMM [Arica,Yarman-Vural  
ICPR00]




Left to right HMM  
[Jelinek, Proc. IEEE 1976]




# One data-driven solution

[Bicego, Cristani, Murino, ICIAP07]

*Sparse HMM*: a HMM with a sparse topology (irrelevant or redundant components are *exactly* 0)



Fully connected model: all transitions are present



Sparse model: many transition probabilities are zero (no connections)

# Sparse HMM

Sparseness is highly desirable:

- It produces a structural simplification of the model, disregarding unimportant parameters
- A sparse model distills the information of all the training data providing only high representative parameters.
- Sparseness is related to generalization ability (Support Vector Machines)

# Some open issues/research trends

## 1. Model selection

- how many states?
- which topology?

## 2. Extending standard models

- modifying dependencies or components

## 3. Injecting discriminative skills into HMM

# Extending standard models (1)

First extension:



adding novel dependencies between components, in order to model different behaviours

Examples:


- Input/Output HMM
- Factorial HMM
- Coupled HMM
- ...

# Preliminary note: the Bayesian Network formalism

Bayes Net: graph where nodes represent variables and edges represent causality




Input-Output HMM: HMM where transitions and emissions are conditional on another sequence (the input sequence)




EX.: finance, the input sequence is a leading market index

# Factorial HMM: more than one state-chain influencing the output



Ex.: speech recognition, where time series generated from several independent sources.

# Coupled HMMs: two interacting HMMs



Ex.: video surveillance, for modelling complex actions like interacting processes



# Extending standard models (2)

Second extension:

- employing as emission probabilities (namely functions modelling output symbols) complex and effective techniques (classifier, distributions,...)

Examples:

- Neural Networks  
[Bourlard, Wellekens, TPAMI 90],...
- Another HMM (to compose Hierarchical HMMs)  
[Fine, Singer, Tishby, ML 98] [Bicego, Grosso, Tistarelli, IVC 09]
- Kernel Machines, such as SVM
- Factor analysis  
[Rosti, Gales, ICASSP 02]
- Generalized Gaussian Distributions  
[Bicego, Gonzalez-Jimenez, Alba-Castro, Grosso, ICPR 08]
- ...

# Extending standard models (2)


- Problems to be faced:
  - full integration of the training of each technique inside the HMM framework
    - “naive” solution: segment data and train separately emissions and other parameters
    - challenging solution: simultaneous training of all parameters
  - in case of Neural Networks or Kernel Machines, it is needed to cast the output of the classifier into a probability value

# HMM application

## 2D shape classification

# 2D shape classification

- Addressed topic in Computer Vision, often basic for three dimensional object recognition
- Fundamental: contour representation
  - Fourier Descriptor
  - chain code
  - curvature based techniques
  - invariants
  - auto-regressive coefficients
  - Hough-based transforms
  - associative memories





# Motivations

- The use of HMM for shape analysis is very poorly addressed
- Previous works:
  - He Kundu (PAMI - 91) using AR coefficients
  - Fred Marques Jorge 1997 (ICIP 97) using chain code
  - Arica Yarman Vural (ICPR 2000) using circular HMM
- Very low entity occlusion
- Closed contours
- Noise sensitivity not analysed

# Objectives

- Investigate the capability of HMM in discriminating object classes, with respect to object translation, rotation, occlusion, noise, and affine projections.
- We use curvature representation for object contour.
- No assumption about HMM topologies or closeness of boundaries.

# Curvature representation



# Curvature representation

- Advantages

- invariant to object translation
- rotation of object is equal to phase translation of the curvature signal;
- can be calculated for open contours

- Disadvantages

- noise sensitivity




# Hidden Markov Model

- Use of Continuous Hidden Markov Model: the emission probability of each state is a Gaussian distribution
- Crucial Issues:
  - Initialisation of training algorithm
  - Model Selection

# HMM Initialisation

- Gaussian Mixture Model clustering of the curvature coefficients: each cluster centroid is used for initialising the parameters of each state.



# HMM model selection

- Bayesian Information Criterion on the initialization
  - 1 HMM model per shape
  - Using BIC on the Gaussian mixture model clustering in order to choose the optimal number of states
  - Advantage: only one HMM training session


# Strategy

- Training: for any object we perform these steps
  - extract edges with Canny edge detector
  - calculate the related curvature signature;
  - train an HMM on it:
    - the HMM was initialised with GMM clustering;
    - the number of HMM states is estimated using the BIC criterion;
    - each HMM was trained using Baum-Welch algorithm
  - at the end of training session we have one HMM  $\lambda_i$  for each object.


## Strategy (cont.)

- Classification: given an unknown sequence  $O$ 
  - compute, for each model  $\lambda_i$ , the probability  $P(O | \lambda_i)$  of generating the sequence  $O$
  - classify  $O$  as belonging to the class whose model shows the highest probability  $P(O | \lambda_i)$ .


## Experimental: The test set



# The models




# The models (2)






# Rotations

- Test set is obtained by rotating 10 times each object by a random angle from 0 to  $2\pi$ .
- Results: Accuracy 100%



# Occlusions

- Each object is occluded: occlusion vary from 5% to 50% (only an half of the whole object is visible)





# Occlusions: results

<u>Occlusion percentage level</u>	<u>Classification Accuracy</u>
5%	100%
10%	100%
15%	100%
20%	100%
25%	100%
30%	100%
35%	100%
40%	97.5%
45%	96.25%
50%	95%


# Noise

- A Gaussian Noise (with mean 0 and variance  $\sigma^2$ ) is added to the X Y coordinates of the object
- $\sigma^2$  varies from 1 to 5: Accuracy 100%. The gaussian filter applied before calculating the curvature is able to remove completely this kind of noise




# Alternative Noise


- Adding noise to the first derivative




$\sigma^2=0.3$



$\sigma^2=0.5$



$\sigma^2=0.7$



$\sigma^2=0.9$

## Noise: results

Noise variance $\sigma^2$	Classification Accuracy
0.1	100.00%
0.3	97.50%
0.5	88.75%
0.7	82.50%
0.9	73.75%

# Occlusions and Rotations: results


























<u>Occlusion percentage level</u>	<u>Classification Accuracy</u>
5%	100%
10%	100%
15%	100%
20%	100%
25%	96.25%
30%	96.25%
35%	95%
40%	91.25%
45%	85%
50%	87.5%

# Occlusions, Rotations and Noise: Results

Occlusion Percentage level	Classification Accuracy		
	Noise $\sigma^2=0.1$	Noise $\sigma^2=0.3$	Noise $\sigma^2=0.5$
50%	86.25%	83.75%	75.00%
40%	93.75%	87.50%	77.50%
30%	98.75%	90.00%	80.00%
20%	98.75%	93.75%	80.00%
10%	100.00%	97.50%	87.50%



# Slant and Tilt Projections

Angoli proiezione	Tilt = 10	Tilt = 20	Tilt = 30	Tilt = 40	Tilt = 50
Slant = 10					
Slant = 20					
Slant = 30					
Slant = 40					
Slant = 50					

# Slant and Tilt Projections: results

Angoli proiezione	Tilt = 10	Tilt = 20	Tilt = 30	Tilt = 40	Tilt = 50
Slant = 10	8/8	8/8	8/8	7/8	4/8
Slant = 20	8/8	8/8	8/8	7/8	4/8
Slant = 30	8/8	8/8	8/8	7/8	4/8
Slant = 40	8/8	8/8	7/8	5/8	4/8
Slant = 50	8/8	8/8	6/8	4/8	4/8

# Conclusions

- System is able to recognize object that could be translated, rotated and occluded, also in presence of noise.
- Translation invariance: due to Curvature
- Rotation invariance: due to Curvature and HMM
- Occlusion invariance: due to HMM
- Robustness to noise: due to HMM

# HMM application


Video Analysis

# Use of the HMMs: main idea

- Each pixel (signal)  $v$  of the sequence is modeled with an HMM  $\lambda_v = (A, B, \pi)$
- $B = \{\mu_i, \sigma_i^2\}$  represents gray level ranges assumed by the  $v$ -th pixel signal, and

$$b_i(O_v) = N(O_v; \mu_i, \sigma_i^2)$$

- The larger the  $\sigma_i^2$ , the more irregular the corresponding signal
- $A :=$  Markov chain that mirrors the evolution of the gray levels



# The idea

- Define the distances between locations on the basis of the distances between the trained Hidden Markov Models
- The segmentation process is obtained using a spatial clustering of HMMs
- We need to define a similarity measure
  - decide when a group (at least, a couple) of neighboring pixels must be labelled as belonging to the same region
- Using this measure the segmentation is obtained as a standard region growing algorithm

# The similarity measure

- The used similarity measure is:

$$D(i, j) = \frac{1}{2} \left\{ \frac{L_{ij} - L_{jj}}{L_{jj}} + \frac{L_{ji} - L_{ii}}{L_{ii}} \right\}$$

where

$$L_{ij} = P(O_i | \lambda_j)$$

- We use a similar distance, *more robust*, which weighs more the states in which the model stands more time

# Results (real)



[Corridoio.avi](#)

**Image based  
segmentation**



**HMM based  
segmentation**

